

# Data Cleaning

Andrew Jaffe

June 15, 2016

# Data

- ▶ We will be using multiple data sets in this lecture:
  - ▶ Salary, Monument, Circulator, and Restaurant from OpenBaltimore: <https://data.baltimorecity.gov/browse?limitTo=datasets>
  - ▶ Gap Minder - very interesting way of viewing longitudinal data
    - ▶ Data is here - <http://www.gapminder.org/data/>
  - ▶ [http://spreadsheets.google.com/pub?key=rMsQHawT0bBb6\\_U2ESjKXYw&output=xls](http://spreadsheets.google.com/pub?key=rMsQHawT0bBb6_U2ESjKXYw&output=xls)

# Data Cleaning

In general, data cleaning is a process of investigating your data for inaccuracies, or recoding it in a way that makes it more manageable.

**MOST IMPORTANT RULE - LOOK AT YOUR DATA!**

Again - `table`, `summarize`, `is.na`, `any`, `all` are useful.

## Dealing with Missing Data

# Missing data types

One of the most important aspects of data cleaning is missing values.

Types of “missing” data:

- ▶ NA - general missing data
- ▶ NaN - stands for “**N**ot **a** **N**umber”, happens when you do  $0/0$ .
- ▶ Inf and -Inf - Infinity, happens when you take a positive number (or negative number) by 0.

# Finding Missing data

Each missing data type has a function that returns TRUE if the data is missing:

- ▶ NA - `is.na`
- ▶ NaN - `is.nan`
- ▶ Inf and -Inf - `is.infinite`
- ▶ `is.finite` returns FALSE for all missing data and TRUE for non-missing
- ▶ `complete.cases` on a `data.frame`/`matrix` returns TRUE if all values in that row of the object are not missing.

# Missing Data with Logicals

One important aspect (esp with subsetting) is that logical operations return NA for NA values. Think about it, the data could be  $> 2$  or not we don't know, so R says there is no TRUE or FALSE, so that is missing:

```
x = c(0, NA, 2, 3, 4)
x > 2
```

```
[1] FALSE      NA FALSE  TRUE  TRUE
```

# Missing Data with Logicals

What to do? What if we want if  $x > 2$  and  $x$  isn't NA?

Don't do  $x \neq \text{NA}$ , do  $x > 2$  and  $x$  is NOT NA:

```
x != NA
```

```
[1] NA NA NA NA NA
```

```
x > 2 & !is.na(x)
```

```
[1] FALSE FALSE FALSE  TRUE  TRUE
```



# Missing Data with Logicals

What about seeing if a value is equal to multiple values? You can do `(x == 1 | x == 2) & !is.na(x)`, but that is not efficient. Introduce the `%in%` operator:

```
(x == 0 | x == 2) # has NA
```

```
[1] TRUE NA TRUE FALSE FALSE
```

```
(x == 0 | x == 2) & !is.na(x) # No NA
```

```
[1] TRUE FALSE TRUE FALSE FALSE
```

```
x %in% c(0, 2) # NEVER has NA and returns logical
```

```
[1] TRUE FALSE TRUE FALSE FALSE
```

# Missing Data with Operations

Similarly with logicals, operations/arithmetic with NA will result in NAs:

```
x + 2
```

```
[1]  2 NA  4  5  6
```

```
x * 2
```

```
[1]  0 NA  4  6  8
```

# Tables and Tabulations

# Creating One-way Tables

Here we will use `table` to make tabulations of the data. Look at `?table` to see options for missing data.

```
table(x)
```

```
x
0 2 3 4
1 1 1 1
```

```
table(x, useNA = "ifany")
```

```
x
    0     2     3     4 <NA>
1    1    1    1    1     1
```

# Creating One-way Tables

You can set `useNA = "always"` to have it always have a column for NA

```
table(c(0, 1, 2, 3, 2, 3, 3, 2,2, 3),  
      useNA = "always")
```

0	1	2	3	<NA>
1	1	4	4	0

# Creating Two-way Tables

A two-way table. If you pass in 2 vectors, `table` creates a 2-dimensional table.

```
tab <- table(c(0, 1, 2, 3, 2, 3, 3, 2, 2, 3),  
             c(0, 1, 2, 3, 2, 3, 3, 4, 4, 3),  
             useNA = "always")
```

# Finding Row or Column Totals

`margin.table` finds the marginal sums of the table. `margin` is 1 for rows, 2 for columns in general in R. Here is the column sums of the table:

```
margin.table(tab, 2)
```

0	1	2	3	4	<NA>
1	1	2	4	2	0

## Proportion Tables

`prop.table` finds the marginal proportions of the table. Think of it dividing the table by it's respective marginal totals. If `margin` not set, divides by overall total.

```
prop.table(tab)
```

	0	1	2	3	4	<NA>
0	0.1	0.0	0.0	0.0	0.0	0.0
1	0.0	0.1	0.0	0.0	0.0	0.0
2	0.0	0.0	0.2	0.0	0.2	0.0
3	0.0	0.0	0.0	0.4	0.0	0.0
<NA>	0.0	0.0	0.0	0.0	0.0	0.0

```
prop.table(tab,1)
```

	0	1	2	3	4	<NA>
0	1.0	0.0	0.0	0.0	0.0	0.0



# Download Salary FY2014 Data

From <https://data.baltimorecity.gov/City-Government/Baltimore-City-Employee-Salaries-FY2014/2j28-xzd7>  
[http://www.aejaffe.com/summerR\\_2016/data/Baltimore\\_City\\_Employee\\_Salaries\\_FY2014.csv](http://www.aejaffe.com/summerR_2016/data/Baltimore_City_Employee_Salaries_FY2014.csv)  
Read the CSV into R Sal:

```
Sal = read.csv("http://www.aejaffe.com/summerR_2016/data/Baltimore_City_Employee_Salaries_FY2014.csv",  
               as.is = TRUE)
```

## Checking for logical conditions

- ▶ `any()` - checks if there are any TRUEs
- ▶ `all()` - checks if ALL are true

```
head(Sal,2)
```

	Name	JobTitle	AgencyID
1	Aaron,Keontae E	AIDE BLUE CHIP	W02200
2	Aaron,Patricia G	Facilities/Office Services II	A03031

  

	Agency	HireDate	AnnualSalary	GrossPay
1	Youth Summer	06/10/2013	\$11310.00	\$873.63
2	OED-Employment Dev	10/24/1979	\$53428.00	\$52868.38

```
any(is.na(Sal$Name)) # are there any NAs?
```

```
[1] FALSE
```



## Example of Recoding: base R

For example, let's say gender was coded as Male, M, m, Female, F, f. Using Excel to find all of these would be a matter of filtering and changing all by hand or using if statements.

In R, you can simply do something like:

```
data$gender[data$gender %in%  
  c("Male", "M", "m")] <- "Male"
```

## Example of Recoding with recode: car package

You can also recode a vector:

```
library(car, quietly = TRUE)
x = rep(c("Male", "M", "m", "f", "Female", "female" ),
        each = 3)
car::recode(x, "c('m', 'M', 'male') = 'Male';
              c('f', 'F', 'female') = 'Female';")
```

```
[1] "Male"    "Male"    "Male"    "Male"    "Male"    "Male"
[8] "Male"    "Male"    "Female"  "Female"  "Female"  "Female"
[15] "Female"  "Female"  "Female"  "Female"
```

## Example of Recoding with revalue: plyr

You can also revalue a vector with the revalue command

```
library(plyr)
plyr::revalue(x, c("M" = "Male", "m" = "Male",
                   "f" = "Female", "female" = "Female"))
```

```
[1] "Male"    "Male"    "Male"    "Male"    "Male"    "Male"
[8] "Male"    "Male"    "Female"  "Female"  "Female"  "Female"
[15] "Female"  "Female"  "Female"  "Female"
```

## Example of Cleaning: more complicated

Sometimes though, it's not so simple. That's where functions that find patterns come in very useful.

```
table(gender)
```

gender

F	FeMAle	FEMALE	Fm	M	Ma	mAle	Male	M
75	82	74	89	89	79	87	89	
Man	Woman							
73	80							

# String functions



## Pasting strings with paste and paste0

Paste can be very useful for joining vectors together:

```
paste("Visit", 1:5, sep = "_")
```

```
[1] "Visit_1" "Visit_2" "Visit_3" "Visit_4" "Visit_5"
```

```
paste("Visit", 1:5, sep = "_", collapse = " ")
```

```
[1] "Visit_1 Visit_2 Visit_3 Visit_4 Visit_5"
```

```
paste("To", "is going be the ", "we go to the store!", sep
```

```
[1] "Today is going be the day we go to the store!"
```

*# and paste0 can be even simpler see ?paste0*

```
paste0("Visit",1:5)
```

```
[1] "Visit1" "Visit2" "Visit3" "Visit4" "Visit5"
```

## Paste Depicting How Collapse Works

```
paste(1:5)
```

```
[1] "1" "2" "3" "4" "5"
```

```
paste(1:5, collapse = " ")
```

```
[1] "1 2 3 4 5"
```

# Useful String Functions

## Useful String functions

- ▶ `toupper()`, `tolower()` - uppercase or lowercase your data:
- ▶ `str_trim()` (in the `stringr` package) or `trimws` in base
  - ▶ will trim whitespace
- ▶ `nchar` - get the number of characters in a string
- ▶ `paste()` - paste strings together with a space
- ▶ `paste0` - paste strings together with no space as default

# The stringr package

Like dplyr, the stringr package:

- ▶ Makes some things more intuitive
- ▶ Is different than base R
- ▶ Is used on forums for answers
- ▶ Has a standard format for most functions
  - ▶ the first argument is a string like first argument is a `data.frame` in dplyr

# Splitting/Find/Replace and Regular Expressions

- ▶ R can do much more than find exact matches for a whole string
- ▶ Like Perl and other languages, it can use regular expressions.
- ▶ What are regular expressions?
  - ▶ Ways to search for specific strings
  - ▶ Can be very complicated or simple
  - ▶ Highly Useful - think “Find” on steroids

# A bit on Regular Expressions

- ▶ `http:`  
`//www.regular-expressions.info/reference.html`
- ▶ They can use to match a large number of strings in one statement
- ▶ `.` matches any single character
- ▶ `*` means repeat as many (even if 0) more times the last character
- ▶ `?` makes the last thing optional
- ▶ `^` matches start of vector `^a` - starts with "a"
- ▶ `$` matches end of vector `b$` - ends with "b"

# Splitting Strings

# Substringing

Very similar:

Base R

- ▶ `substr(x, start, stop)` - substrings from position start to position stop
- ▶ `strsplit(x, split)` - splits strings up - returns list!

stringr

- ▶ `str_sub(x, start, end)` - substrings from position start to position end
- ▶ `str_split(string, pattern)` - splits strings up - returns list!



## Splitting String: base R

In base R, `strsplit` splits a vector on a string into a list

```
x <- c("I really", "like writing", "R code programs")  
y <- strsplit(x, split = " ") # returns a list  
y
```

```
[[1]]
```

```
[1] "I"      "really"
```

```
[[2]]
```

```
[1] "like"   "writing"
```

```
[[3]]
```

```
[1] "R"      "code"   "programs"
```

## Splitting String: stringr

stringr::str\_split do the same thing:

```
library(stringr)
y2 <- str_split(x, " ") # returns a list
y2
```

```
[[1]]
```

```
[1] "I"      "really"
```

```
[[2]]
```

```
[1] "like"    "writing"
```

```
[[3]]
```

```
[1] "R"      "code"    "programs"
```

## Using a fixed expression

One example case is when you want to split on a period “.”. In regular expressions `.` means **ANY** character, so

```
str_split("I.like.strings", ".")
```

```
[[1]]  
[1] "" "" "" "" "" "" "" "" "" "" "" "" "" "" "" "" "" "" ""
```

```
str_split("I.like.strings", fixed("."))
```

```
[[1]]  
[1] "I"      "like"    "strings"
```

## Let's extract from y

```
suppressPackageStartupMessages(library(dplyr)) # must be loaded  
y[[2]]
```

```
[1] "like"      "writing"
```

```
sapply(y, dplyr::first) # on the fly
```

```
[1] "I"        "like" "R"
```

```
sapply(y, nth, 2) # on the fly
```

```
[1] "really"    "writing" "code"
```

```
sapply(y, last) # on the fly
```

```
[1] "really"    "writing"  "programs"
```

## 'Find' functions: base R

`grep`: `grep`, `grep1`, `regexpr` and `gregexpr` search for matches to argument `pattern` within each element of a character vector: they differ in the format of and amount of detail in the results.

`grep(pattern, x, fixed=FALSE)`, where:

- ▶ `pattern` = character string containing a regular expression to be matched in the given character vector.
- ▶ `x` = a character vector where matches are sought, or an object which can be coerced by `as.character` to a character vector.
- ▶ If `fixed=TRUE`, it will do exact matching for the phrase anywhere in the vector (regular find)

## 'Find' functions: `stringr`

`str_detect`, `str_subset`, `str_replace`, and `str_replace_all` search for matches to argument `pattern` within each element of a character vector: they differ in the format of and amount of detail in the results.

- ▶ `str_detect` - returns `TRUE` if `pattern` is found
- ▶ `str_subset` - returns only the strings which `pattern` were detected
  - ▶ convenient wrapper around `x[str_detect(x, pattern)]`
- ▶ `str_extract` - returns only strings which `pattern` were detected, but **ONLY** the `pattern`
- ▶ `str_replace` - replaces `pattern` with `replacement` the first time
- ▶ `str_replace_all` - replaces `pattern` with `replacement` as many times matched

## 'Find' functions: stringr compared to base R

Base R does not use these functions. Here is a “translator” of the `stringr` function to base R functions

- ▶ `str_detect` - similar to `grepl` (return logical)
- ▶ `grep(value = FALSE)` is similar to `which(str_detect())`
- ▶ `str_subset` - similar to `grep(value = TRUE)` - return value of matched
- ▶ `str_replace` - similar to `sub` - replace one time
- ▶ `str_replace_all` - similar to `gsub` - replace many times

# Let's look at modifier for stringr

?modifiers

- ▶ `fixed` - match everything exactly
- ▶ `regex` - default - uses **regular expressions**
- ▶ `ignore_case` is an option to not have to use `tolower`



# Important Comparisons

Base R:

- ▶ Argument order is (pattern, x)
- ▶ Uses option (fixed = TRUE)

stringr

- ▶ Argument order is (string, pattern) aka (x, pattern)
- ▶ Uses function fixed(pattern)

## 'Find' functions: Finding Indices

These are the indices where the pattern match occurs:

```
grep("Rawlings", Sal$Name)
```

```
[1] 13832 13833 13834 13835
```

```
which(grepl("Rawlings", Sal$Name))
```

```
[1] 13832 13833 13834 13835
```

```
which(str_detect(Sal$Name, "Rawlings"))
```

```
[1] 13832 13833 13834 13835
```

## 'Find' functions: Finding Logicals

These are the indices where the pattern match occurs:

```
head(grepl("Rawlings",Sal$Name))
```

```
[1] FALSE FALSE FALSE FALSE FALSE FALSE
```

```
head(str_detect(Sal$Name, "Rawlings"))
```

```
[1] FALSE FALSE FALSE FALSE FALSE FALSE
```

## 'Find' functions: finding values, base R

```
grep("Rawlings",Sal$Name,value=TRUE)
```

```
[1] "Rawlings,Kellye A"           "Rawlings,MarqWell D"  
[3] "Rawlings,Paula M"           "Rawlings-Blake,Stephanie C"
```

```
Sal[grep("Rawlings",Sal$Name),]
```

		Name	JobTitle	Agency
13832		Rawlings,Kellye A	EMERGENCY DISPATCHER	A40
13833		Rawlings,MarqWell D	AIDE BLUE CHIP	W02
13834		Rawlings,Paula M	COMMUNITY AIDE	A04
13835		Rawlings-Blake,Stephanie C	MAYOR	A01

  

		Agency	HireDate	AnnualSalary	GrossPay
13832	M-R	Info Technology	01/06/2003	\$47980.00	\$68426.00
13833		Youth Summer	06/15/2012	\$11310.00	\$507.00
13834		R&P-Recreation	12/10/2007	\$19802.00	\$8195.00
13835		Mayors Office	12/07/1995	\$163365.00	\$161219.00

## 'Find' functions: finding values, stringr and dplyr

```
str_subset(Sal$Name, "Rawlings")
```

```
[1] "Rawlings,Kellye A"          "Rawlings,MarqWell D"
[3] "Rawlings,Paula M"          "Rawlings-Blake,Stephanie
```

```
Sal %>% filter(str_detect(Name, "Rawlings"))
```

	Name	JobTitle	AgencyID
1	Rawlings,Kellye A	EMERGENCY DISPATCHER	A40302
2	Rawlings,MarqWell D	AIDE BLUE CHIP	W02384
3	Rawlings,Paula M	COMMUNITY AIDE	A04015
4	Rawlings-Blake,Stephanie C	MAYOR	A01001

  

	Agency	HireDate	AnnualSalary	GrossPay
1	M-R Info Technology	01/06/2003	\$47980.00	\$68426.73
2	Youth Summer	06/15/2012	\$11310.00	\$507.50
3	R&P-Recreation	12/10/2007	\$19802.00	\$8195.79
4	Mayors Office	12/07/1995	\$163365.00	\$161219.24

## Showing difference in str\_extract

str\_extract extracts just the matched string

```
ss = str_extract(Sal$Name, "Rawling")  
head(ss)
```

```
[1] NA NA NA NA NA NA
```

```
ss[ !is.na(ss)]
```

```
[1] "Rawling" "Rawling" "Rawling" "Rawling"
```

## Showing difference in str\_extract and str\_extract\_all

str\_extract\_all extracts all the matched strings

```
head(str_extract(Sal$AgencyID, "\\d"))
```

```
[1] "0" "0" "2" "6" "9" "4"
```

```
head(str_extract_all(Sal$AgencyID, "\\d"), 2)
```

```
[[1]]
```

```
[1] "0" "2" "2" "0" "0"
```

```
[[2]]
```

```
[1] "0" "3" "0" "3" "1"
```

# Using Regular Expressions

- ▶ Look for any name that starts with:
  - ▶ Payne at the beginning,
  - ▶ Leonard and then an S
  - ▶ Spence then capital C

```
head(grep("^Payne.*", x = Sal$Name, value = TRUE), 3)
```

```
[1] "Payne El,Jackie"           "Payne Johnson,Nickole A"  
[3] "Payne,Chanel"
```

```
head(grep("Leonard.?S", x = Sal$Name, value = TRUE))
```

```
[1] "Payne,Leonard S"          "Szumlanski,Leonard S"
```

```
head(grep("Spence.*C.*", x = Sal$Name, value = TRUE))
```

```
[1] "Greene,Spencer C"         "Spencer,Charles A"      "Spencer,Clarence W"  
[4] "Spencer,Clarence W"      "Spencer,Michael C"
```



## Using Regular Expressions: stringr

```
head(str_subset( Sal$Name, "^Payne.*"), 3)
```

```
[1] "Payne El,Jackie"          "Payne Johnson,Nickole A"  
[3] "Payne,Chanel"
```

```
head(str_subset( Sal$Name, "Leonard.?S"))
```

```
[1] "Payne,Leonard S"          "Szumlanski,Leonard S"
```

```
head(str_subset( Sal$Name, "Spence.*C.*"))
```

```
[1] "Greene,Spencer C"        "Spencer,Charles A"      "Spencer,Clarence W"  
[4] "Spencer,Clarence W"      "Spencer,Michael C"
```

# Replace

Let's say we wanted to sort the data set by Annual Salary:

```
class(Sal$AnnualSalary)
```

```
[1] "character"
```

```
sort(c("1", "2", "10")) # not sort correctly (order simply)
```

```
[1] "1"  "10" "2"
```

```
order(c("1", "2", "10"))
```

```
[1] 1 3 2
```

# Replace

So we must change the annual pay into a numeric:

```
head(Sal$AnnualSalary, 4)
```

```
[1] "$11310.00" "$53428.00" "$68300.00" "$62000.00"
```

```
head(as.numeric(Sal$AnnualSalary), 4)
```

```
Warning in head(as.numeric(Sal$AnnualSalary), 4): NAs introduced by coercion
```

```
[1] NA NA NA NA
```

R didn't like the \$ so it thought turned them all to NA.  
sub() and gsub() can do the replacing part in base R.

## Replacing and subbing

Now we can replace the \$ with nothing (used fixed=TRUE because \$ means ending):

```
Sal$AnnualSalary <- as.numeric(gsub(pattern = "$", replacement = "",  
                                   Sal$AnnualSalary, fixed=TRUE))  
Sal <- Sal[order(Sal$AnnualSalary, decreasing=TRUE), ]  
Sal[1:5, c("Name", "AnnualSalary", "JobTitle")]
```

	Name	AnnualSalary	JobTitle
1222	Bernstein, Gregg L	238772	STATE'S ATTORNEY
3175	Charles, Ronnie E	200000	EXECUTIVE LEVEL III
985	Batts, Anthony W	193800	EXECUTIVE LEVEL III
1343	Black, Harry E	190000	EXECUTIVE LEVEL III
16352	Swift, Michael	187200	CONTRACT SERV SPEC II

## Replacing and subbing: stringr

We can do the same thing (with 2 piping operations!) in dplyr

```
dplyr_sal = Sal
dplyr_sal = dplyr_sal %>% mutate(
  AnnualSalary = AnnualSalary %>%
    str_replace(
      fixed("$"),
      "") %>%
    as.numeric() %>%
    arrange(desc(AnnualSalary))
check_Sal = Sal
rownames(check_Sal) = NULL
all.equal(check_Sal, dplyr_sal)
```

[1] TRUE