Data Summarization

Andrew Jaffe

June 15, 2016

Data Summarization

- Basic statistical summarization
    - `mean(x)`: takes the mean of x
    - `sd(x)`: takes the standard deviation of x
    - `median(x)`: takes the median of x
    - `quantile(x)`: displays sample quantities of x. Default is min, IQR, max
    - `range(x)`: displays the range. Same as c(min(x), max(x))

Some examples

We can use the `mtcars` and Charm City Circulator datasets to explore different ways of summarizing data.

```
head(mtcars)
```

```
##                    mpg cyl disp  hp drat    wt  qsec vs
## Mazda RX4         21.0   6  160 110 3.90 2.620 16.46  0
## Mazda RX4 Wag     21.0   6  160 110 3.90 2.875 17.02  0
## Datsun 710        22.8   4  108  93 3.85 2.320 18.61  1
## Hornet 4 Drive    21.4   6  258 110 3.08 3.215 19.44  1
## Hornet Sportabout 18.7   8  360 175 3.15 3.440 17.02  0
## Valiant           18.1   6  225 105 2.76 3.460 20.22  1
```

Statistical summarization

```r
mean(mtcars$hp)
```

```
## [1] 146.6875
```

```r
quantile(mtcars$hp)
```

```
##    0%   25%   50%   75%  100%
##  52.0  96.5 123.0 180.0 335.0
```

Statistical summarization

```r
median(mtcars$wt)
```

```
## [1] 3.325
```

```r
quantile(mtcars$wt, probs = 0.6)
```

```
##  60%
## 3.44
```

Statistical summarization

Note that many of these functions have additional inputs regarding missing data, typically requiring the `na.rm` argument.

```r
x = c(1,5,7,NA,4,2, 8,10,45,42)
mean(x)
```

```
## [1] NA
```

```r
mean(x,na.rm=TRUE)
```

```
## [1] 13.77778
```

```r
quantile(x,na.rm=TRUE)
```

```
##   0%  25%  50%  75% 100%
##    1    4    7   10   45
```

Data Summarization on matrices/data frames

- Basic statistical summarization
    - `rowMeans(x)`: takes the means of each row of x
    - `colMeans(x)`: takes the means of each column of x
    - `rowSums(x)`: takes the sum of each row of x
    - `colSums(x)`: takes the sum of each column of x
    - `summary(x)`: for data frames, displays the quantile information

Charm City Circulator data

Please download the Charm City Circulator data:
http://www.aejaffe.com/summerR_2016/data/Charm_City_
Circulator_Ridership.csv

```
circ = read.csv("http://www.aejaffe.com/summerR_2016/data/C
                header=TRUE,as.is=TRUE)
```

Subsetting to specific columns

Let's just take columns that represent average ridership:

```r
library(dplyr,quietly = TRUE)
```

```
##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##     filter, lag

## The following objects are masked from 'package:base':
##
##     intersect, setdiff, setequal, union
```

```r
circ2 = select(circ, date, day, ends_with("Average"))
```

column and row means

```r
avgs = select(circ2, ends_with("Average"))
colMeans(avgs,na.rm=TRUE)
```

```
## orangeAverage purpleAverage   greenAverage bannerAverage
##      3033.1611     4016.9345      1957.7814      827.2685
```

```r
circ2$daily = rowMeans(avgs,na.rm=TRUE)
head(circ2$daily)
```

```
## [1]   952.0   796.0 1211.5 1213.5 1644.0 1490.5
```

Summary

```
summary(circ2)
```

```
##     date              day              orangeAverage   pu
## Length:1146       Length:1146        Min.   :   0     Mi
## Class :character  Class :character   1st Qu.:2001     1s
## Mode  :character  Mode  :character   Median :2968     Me
##                                      Mean   :3033     Me
##                                      3rd Qu.:4020     3r
##                                      Max.   :6926     Ma
##                                      NA's   :10       NA
##   greenAverage   bannerAverage        daily
## Min.   :   0    Min.   :   0.0    Min.   :   0
## 1st Qu.:1491    1st Qu.: 632.5    1st Qu.:2097
## Median :2079    Median : 763.0    Median :2846
## Mean   :1958    Mean   : 827.3    Mean   :2878
## 3rd Qu.:2340    3rd Qu.: 945.9    3rd Qu.:3646
## Max.   :5094    Max.   :4617.0    Max.   :6123
## NA's   :661     NA's   :876       NA's   :10
```

Apply statements

You can apply more general functions to the rows or columns of a matrix or data frame, beyond the mean and sum.

```
apply(X, MARGIN, FUN, ...)
```

> *X : an array, including a matrix.*
> *MARGIN : a vector giving the subscripts which the function will be applied over. E.g., for a matrix 1 indicates rows, 2 indicates columns, c(1, 2) indicates rows and columns. Where X has named dimnames, it can be a character vector selecting dimension names.*
> *FUN : the function to be applied: see 'Details'.*
> *... : optional arguments to FUN.*

Apply statements

```
apply(avgs,2,mean,na.rm=TRUE) # column means
```

```
## orangeAverage purpleAverage   greenAverage bannerAverage
##      3033.1611      4016.9345      1957.7814      827.2685
```

```
apply(avgs,2,sd,na.rm=TRUE) # columns sds
```

```
## orangeAverage purpleAverage   greenAverage bannerAverage
##      1227.5779      1406.6544       592.8969      436.0487
```

```
apply(avgs,2,max,na.rm=TRUE) # column maxs
```

```
## orangeAverage purpleAverage   greenAverage bannerAverage
##         6926.5         8089.5         5094.0         4617.0
```

Other Apply Statements

- `tapply()`: 'table' apply
- `lapply()`: 'list' apply [tomorrow]
- `sapply()`: 'simple' apply [tomorrow]
- Other less used ones. . .

See more details here: `http://nsaunders.wordpress.com/2010/08/20/a-brief-introduction-to-apply-in-r/`

tapply()

From the help file: "Apply a function to each cell of a ragged array, that is to each (non-empty) group of values given by a unique combination of the levels of certain factors."

```
tapply(X, INDEX, FUN = NULL, ..., simplify = TRUE)
```

Simply put, you can apply function FUN to X within each categorical level of INDEX. It is very useful for assessing properties of continuous data by levels of categorical data.

tapply()

For example, we can estimate the highest average daily ridership for each day of the week in 1 line in the Circulator dataset.

```
tapply(circ2$daily, circ2$day, max, na.rm=TRUE)
```

```
##    Friday    Monday  Saturday    Sunday  Thursday    Tues
##   5600.75   5002.25   6123.00   3980.25   4820.50   4855
```

Data Summarization

- Basic summarization plots
  - `plot(x,y)`: scatterplot of x and y
  - `boxplot(y~x)`: boxplot of y against levels of x
  - `hist(x)`: histogram of x
  - `density(x)`: kernel density plot of x

Basic Plots

Plotting is an important component of exploratory data analysis.
We will review some of the more useful and informative plots here.
We will go over formatting and making plots look nicer in additional
lectures.

Scatterplot

```
plot(mtcars$mpg, mtcars$disp)
```

Histograms

```
hist(circ2$daily)
```


Histogram of circ2$daily

Density

```
## plot(density(circ2$daily))
plot(density(circ2$daily,na.rm=TRUE))
```

density.default(x = circ2$daily, na.rm = TRUE)

N = 1136   Bandwidth = 232.1

Boxplots

```
boxplot(circ2$daily ~ circ2$day)
```

Boxplots

```
boxplot(daily ~ day, data=circ2)
```

Data Summarization for data.frames

- Basic summarization plots
  - `matplot(x,y)`: scatterplot of two matrices, x and y
  - `pairs(x,y)`: plots pairwise scatter plots of matrices x and y, column by column

Matrix plot

```
matplot(avgs)
```