

Manipulating Data in R

John Muschelli

June 15, 2016

Overview

In this module, we will show you how to:

1. Reshaping data from long (tall) to wide (fat)
2. Reshaping data from wide (fat) to long (tall)
3. Merging Data
4. Perform operations by a grouping variable

Setup

We will show you how to do each operation in base R then show you how to use the `dplyr` or `tidyr` package to do the same operation (if applicable).

See the “Data Wrangling Cheat Sheet using `dplyr` and `tidyr`”:

- ▶ <https://www.rstudio.com/wp-content/uploads/2015/02/data-wrangling-cheatsheet.pdf>

Load the packages/libraries

```
library(dplyr)  
library(tidyr)
```

Data used: Charm City Circulator

`http://www.aejaffe.com/winterR_2016/data/Charm_City_Circulator_Ridership.csv`

Let's read in the Charm City Circulator data:

```
circ = read.csv("http://www.aejaffe.com/winterR_2016/data/CircData.csv")
head(circ, 2)
```

```

      day      date orangeBoardings orangeAlightings orange
1 Monday 01/11/2010             877             1027
2 Tuesday 01/12/2010             777             815
  purpleBoardings purpleAlightings purpleAverage greenBoard
1              NA              NA              NA
2              NA              NA              NA
  greenAlightings greenAverage bannerBoardings bannerAlight
1              NA              NA              NA
2              NA              NA              NA
  bannerAverage daily
1              NA    952
2              NA    796

```

Creating a Date class from a character date

The lubridate package is great for dates:

```
library(lubridate) # great for dates!
```

Attaching package: 'lubridate'

The following object is masked from 'package:base':

date

```
circ = mutate(circ, date = mdy(date))  
sum( is.na(circ$date) )
```

```
[1] 0
```

```
head(circ$date)
```

```
[1] "2010-01-11" "2010-01-12" "2010-01-13" "2010-01-14" "2010-01-15"  
[6] "2010-01-16"
```

Making column names a little more separated

We will use `str_replace` from `stringr` to put periods in the column names.

```
library(stringr)
cn = colnames(circ)
cn = cn %>%
  str_replace("Board", ".Board") %>%
  str_replace("Alight", ".Alight") %>%
  str_replace("Average", ".Average")
colnames(circ) = cn
cn
```

[1]	"day"	"date"	"orange.Boardi
[4]	"orange.Alightings"	"orange.Average"	"purple.Boardi
[7]	"purple.Alightings"	"purple.Average"	"green.Boardin
[10]	"green.Alightings"	"green.Average"	"banner.Boardi
[13]	"banner.Alightings"	"banner.Average"	"daily"

Removing the daily ridership

We want to look at each ridership, and will remove the `daily` column:

```
circ$daily = NULL
```


Reshaping data from wide (fat) to long (tall)

See http://www.cookbook-r.com/Manipulating_data/Converting_data_between_wide_and_long_format/

- ▶ Wide - multiple columns per observation
 - ▶ e.g. visit1, visit2, visit3

	id	visit1	visit2	visit3
1	1	10	4	3
2	2	5	6	NA

- ▶ Long - multiple rows per observation

	id	visit	value
1	1	1	10
2	1	2	4
3	1	3	3
4	2	1	5
5	2	2	6

Reshaping data from wide (fat) to long (tall): base R

The reshape command exists. It is a **confusing** function. Don't use it.

Reshaping data from wide (fat) to long (tall): tidyr

`tidyr::gather` - puts column data into rows.

We want the column names into “var” variable in the output dataset and the value in “number” variable. We then describe which columns we want to “gather:”

```
long = gather(circ, key = "var", value = "number",  
              starts_with("orange"),  
              starts_with("purple"),  
              starts_with("green"),  
              starts_with("banner"))  
head(long, 2)
```

	day	date	var	number
1	Monday	2010-01-11	orange.Boardings	877
2	Tuesday	2010-01-12	orange.Boardings	777

```
table(long$var)
```

Reshaping data from wide (fat) to long (tall): tidy

Now each var is boardings, averages, or alightings. We want to separate these so we can have these by line.

```
long = separate_(long, "var",  
                  into = c("line", "type"),  
                  sep = "[.]")  
  
head(long, 3)
```

	day	date	line	type	number
1	Monday	2010-01-11	orange	Boardings	877
2	Tuesday	2010-01-12	orange	Boardings	777
3	Wednesday	2010-01-13	orange	Boardings	1203

```
unique(long$line)
```

```
[1] "orange" "purple" "green"  "banner"
```

```
unique(long$type)
```

```
[1] "Boardings" "Alightings" "Averages"
```

Reshaping data from long (tall) to wide (fat): tidyr

In tidyr, the spread function spreads rows into columns. Now we have a long data set, but we want to separate the Average, Alightings and Boardings into different columns:

```
# have to remove missing days  
wide = filter(long, !is.na(date))  
wide = spread(wide, type, number)  
head(wide)
```

	day	date	line	Alightings	Average	Boardings
1	Friday	2010-01-15	banner	NA	NA	NA
2	Friday	2010-01-15	green	NA	NA	NA
3	Friday	2010-01-15	orange	1643	1644	1645
4	Friday	2010-01-15	purple	NA	NA	NA
5	Friday	2010-01-22	banner	NA	NA	NA
6	Friday	2010-01-22	green	NA	NA	NA

Reshaping data from long (tall) to wide (fat): tidy

We can use `rowSums` to see if any values in the row is NA and keep if the row, which is a combination of date and line type has any non-missing data.

```
# wide = wide %>%  
#   select(Alightings, Average, Boardings) %>%  
#   mutate(good = rowSums(is.na(.)) > 0)  
namat = !is.na(select(wide, Alightings, Average, Boardings))  
head(namat)
```

	Alightings	Average	Boardings
1	FALSE	FALSE	FALSE
2	FALSE	FALSE	FALSE
3	TRUE	TRUE	TRUE
4	FALSE	FALSE	FALSE
5	FALSE	FALSE	FALSE
6	FALSE	FALSE	FALSE

```
wide$good = rowSums(namat) > 0
```

Reshaping data from long (tall) to wide (fat): tidyr

Now we can filter only the good rows and delete the good column.

```
wide = filter(wide, good) %>% select(-good)
head(wide)
```

	day	date	line	Alightings	Average	Boardings
1	Friday	2010-01-15	orange	1643	1644.0	1645
2	Friday	2010-01-22	orange	1388	1394.5	1401
3	Friday	2010-01-29	orange	1322	1332.0	1342
4	Friday	2010-02-05	orange	1204	1217.5	1231
5	Friday	2010-02-12	orange	678	671.0	664
6	Friday	2010-02-19	orange	1647	1642.0	1637

Data Merging/Append in Base R

- ▶ Merging - joining data sets together - usually on key variables, usually “id”
- ▶ `merge()` is the most common way to do this with data sets
- ▶ `rbind/cbind` - row/column bind, respectively
 - ▶ `rbind` is the equivalent of “appending” in Stata or “setting” in SAS
 - ▶ `cbind` allows you to add columns in addition to the previous ways
- ▶ `t()` is a function that will transpose the data

Merging

```
base <- data.frame(id = 1:10, Age= seq(55,60, length=10))  
base[1:2,]
```

	id	Age
1	1	55.00000
2	2	55.55556

```
visits <- data.frame(id = rep(1:8, 3), visit= rep(1:3, 8),  
                     Outcome = seq(10,50, length=24))  
visits[1:2,]
```

	id	visit	Outcome
1	1	1	10.00000
2	2	2	11.73913

Merging

```
merged.data <- merge(base, visits, by="id")  
merged.data[1:5,]
```

	id	Age	visit	Outcome
1	1	55.00000	1	10.00000
2	1	55.00000	3	23.91304
3	1	55.00000	2	37.82609
4	2	55.55556	2	11.73913
5	2	55.55556	1	25.65217

```
dim(merged.data)
```

```
[1] 24 4
```

Merging

```
all.data <- merge(base, visits, by="id", all=TRUE)
tail(all.data)
```

	id	Age	visit	Outcome
21	7	58.33333	2	48.26087
22	8	58.88889	2	22.17391
23	8	58.88889	1	36.08696
24	8	58.88889	3	50.00000
25	9	59.44444	NA	NA
26	10	60.00000	NA	NA

```
dim(all.data)
```

```
[1] 26 4
```

Joining in dplyr

- ▶ ?join - see different types of joining for dplyr
- ▶ Let's look at <https://www.rstudio.com/wp-content/uploads/2015/02/data-wrangling-cheatsheet.pdf>

Left Join

```
lj = left_join(base, visits)
```

Joining by: "id"

```
dim(lj)
```

```
[1] 26  4
```

```
tail(lj)
```

	id	Age	visit	Outcome
21	7	58.33333	2	48.26087
22	8	58.88889	2	22.17391
23	8	58.88889	1	36.08696
24	8	58.88889	3	50.00000
25	9	59.44444	NA	NA
26	10	60.00000	NA	NA

Right Join

```
rj = right_join(base, visits)
```

Joining by: "id"

```
dim(rj)
```

```
[1] 24  4
```

```
tail(rj)
```

	id	Age	visit	Outcome
19	3	56.11111	1	41.30435
20	4	56.66667	2	43.04348
21	5	57.22222	3	44.78261
22	6	57.77778	1	46.52174
23	7	58.33333	2	48.26087
24	8	58.88889	3	50.00000

Full Join

```
fj = full_join(base, visits)
```

Joining by: "id"

```
dim(fj)
```

```
[1] 26  4
```

```
tail(fj)
```

	id	Age	visit	Outcome
21	7	58.33333	2	48.26087
22	8	58.88889	2	22.17391
23	8	58.88889	1	36.08696
24	8	58.88889	3	50.00000
25	9	59.44444	NA	NA
26	10	60.00000	NA	NA

Perform Operations By Groups: base R

The `tapply` command will take in a vector (`X`), perform a function (`FUN`) over an index (`INDEX`):

```
args(tapply)
```

```
function (X, INDEX, FUN = NULL, ..., simplify = TRUE)  
NULL
```


Perform Operations By Groups: base R

Let's get the mean Average ridership by line:

```
tapply(wide$Average, wide$line, mean, na.rm = TRUE)
```

banner	green	orange	purple
827.2685	1957.7814	3033.1611	4016.9345

Perform Operations By Groups: dplyr

Let's get the mean Average ridership by line We will use `group_by` to group the data by line, then use `summarize` (or `summarise`) to get the mean Average ridership:

```
gb = group_by(wide, line)
summarize(gb, mean_avg = mean(Average))
```

Source: local data frame [4 x 2]

	line (chr)	mean_avg (dbl)
1	banner	827.2685
2	green	1957.7814
3	orange	3033.1611
4	purple	4016.9345

Perform Operations By Groups: dplyr with piping

Using piping, this is:

```
wide %>%  
  group_by(line) %>%  
  summarise(mean_avg = mean(Average))
```

Source: local data frame [4 x 2]

	line (chr)	mean_avg (dbl)
1	banner	827.2685
2	green	1957.7814
3	orange	3033.1611
4	purple	4016.9345

Perform Operations By Multiple Groups: dplyr

This can easily be extended using `group_by` with multiple groups. Let's define the year of riding:

```
wide = wide %>% mutate(year = year(date),  
                        month = month(date))  
  
wide %>%  
  group_by(line, year) %>%  
  summarise(mean_avg = mean(Average))
```

Source: local data frame [13 x 3]

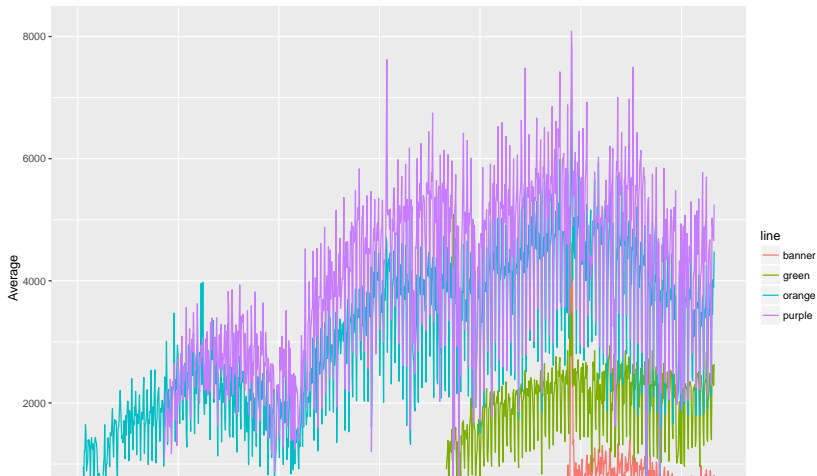
Groups: line [?]

	line (chr)	year (dbl)	mean_avg (dbl)
1	banner	2012	882.0929
2	banner	2013	635.3833
3	green	2011	1455.1667
4	green	2012	2028.7740
5	green	2013	2028.5250

Perform Operations By Multiple Groups: dplyr

We can then easily plot each day over time:

```
library(ggplot2)
ggplot(aes(x = date, y = Average,
           colour = line), data = wide) + geom_line()
```



Perform Operations By Multiple Groups: dplyr

Let's create the middle of the month (the 15th for example), and name it mon.

```
mon = wide %>%  
  dplyr::group_by(line, month, year) %>%  
  dplyr::summarise(mean_avg = mean(Average))  
mon = mutate(mon,  
              mid_month = dmy(paste0("15-", month, "-", year  
head(mon)
```

Source: local data frame [6 x 5]

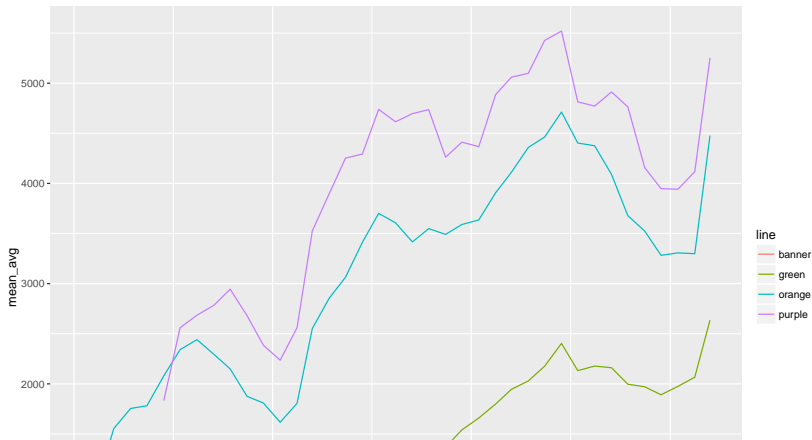
Groups: line, month [6]

	line (chr)	month (dbl)	year (dbl)	mean_avg (dbl)	mid_month (date)
1	banner	1	2013	610.3226	2013-01-15
2	banner	2	2013	656.4643	2013-02-15
3	banner	3	2013	822.0000	2013-03-15
4	banner	6	2012	1288.1296	2012-06-15

Perform Operations By Multiple Groups: dplyr

We can then easily plot the mean of each month to see a smoother output:

```
ggplot(aes(x = mid_month,  
           y = mean_avg,  
           colour = line), data = mon) + geom_line()
```



Bonus! Points with a smoother!

```
ggplot(aes(x = date, y = Average, colour = line),  
       data = wide) + geom_smooth(se = FALSE) +  
       geom_point(size = .5)
```

